

# car-insurance-prediction

April 29, 2023

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
[2]: df=pd.read_csv("D:\\projects\\DS Internship datasets\\project_
↪2\\Car_Insurance_Claim.csv")
```

```
[3]: df.head()
```

```
[3]:      ID    AGE  GENDER    RACE  DRIVING_EXPERIENCE  EDUCATION  \
0  569520   65+  female  majority           0-9y  high school
1  750365  16-25   male  majority           0-9y         none
2  199901  16-25  female  majority           0-9y  high school
3  478866  16-25   male  majority           0-9y  university
4  731664  26-39   male  majority          10-19y         none

      INCOME  CREDIT_SCORE  VEHICLE_OWNERSHIP  VEHICLE_YEAR  MARRIED  \
0  upper class      0.629027           1.0  after 2015      0.0
1  poverty      0.357757           0.0  before 2015      0.0
2  working class      0.493146           1.0  before 2015      0.0
3  working class      0.206013           1.0  before 2015      0.0
4  working class      0.388366           1.0  before 2015      0.0

      CHILDREN  POSTAL_CODE  ANNUAL_MILEAGE  VEHICLE_TYPE  SPEEDING_VIOLATIONS  \
0          1.0         10238        12000.0         sedan              0
1          0.0         10238        16000.0         sedan              0
2          0.0         10238        11000.0         sedan              0
3          1.0         32765        11000.0         sedan              0
4          0.0         32765        12000.0         sedan              2

      DUIS  PAST_ACCIDENTS  OUTCOME
0        0              0        0.0
1        0              0        1.0
2        0              0        0.0
```

3	0	0	0.0
4	0	1	1.0

```
[4]: df.tail()
```

```
[4]:
```

	ID	AGE	GENDER	RACE	DRIVING_EXPERIENCE	EDUCATION	\
9995	323164	26-39	female	majority	10-19y	university	
9996	910346	26-39	female	majority	10-19y	none	
9997	468409	26-39	male	majority	0-9y	high school	
9998	903459	26-39	female	majority	10-19y	high school	
9999	442696	26-39	female	majority	0-9y	none	

	INCOME	CREDIT_SCORE	VEHICLE_OWNERSHIP	VEHICLE_YEAR	MARRIED	\
9995	upper class	0.582787	1.0	before 2015	0.0	
9996	middle class	0.522231	1.0	after 2015	0.0	
9997	middle class	0.470940	1.0	before 2015	0.0	
9998	poverty	0.364185	0.0	before 2015	0.0	
9999	working class	0.435225	1.0	before 2015	1.0	

	CHILDREN	POSTAL_CODE	ANNUAL_MILEAGE	VEHICLE_TYPE	SPEEDING_VIOLATIONS	\
9995	0.0	10238	16000.0	sedan	0	
9996	1.0	32765	NaN	sedan	1	
9997	1.0	10238	14000.0	sedan	0	
9998	1.0	10238	13000.0	sedan	2	
9999	1.0	10238	13000.0	sedan	0	

	DUIS	PAST_ACCIDENTS	OUTCOME
9995	0	1	0.0
9996	0	0	0.0
9997	0	0	0.0
9998	0	1	1.0
9999	0	0	0.0

```
[5]: df.describe()
```

```
[5]:
```

	ID	CREDIT_SCORE	VEHICLE_OWNERSHIP	MARRIED	\
count	10000.000000	9018.000000	10000.000000	10000.000000	
mean	500521.906800	0.515813	0.697000	0.498200	
std	290030.768758	0.137688	0.459578	0.500022	
min	101.000000	0.053358	0.000000	0.000000	
25%	249638.500000	0.417191	0.000000	0.000000	
50%	501777.000000	0.525033	1.000000	0.000000	
75%	753974.500000	0.618312	1.000000	1.000000	
max	999976.000000	0.960819	1.000000	1.000000	

	CHILDREN	POSTAL_CODE	ANNUAL_MILEAGE	SPEEDING_VIOLATIONS	\
count	10000.000000	10000.000000	9043.000000	10000.000000	

mean	0.688800	19864.548400	11697.003207	1.482900
std	0.463008	18915.613855	2818.434528	2.241966
min	0.000000	10238.000000	2000.000000	0.000000
25%	0.000000	10238.000000	10000.000000	0.000000
50%	1.000000	10238.000000	12000.000000	0.000000
75%	1.000000	32765.000000	14000.000000	2.000000
max	1.000000	92101.000000	22000.000000	22.000000

	DUIS	PAST_ACCIDENTS	OUTCOME
count	10000.00000	10000.000000	10000.000000
mean	0.23920	1.056300	0.313300
std	0.55499	1.652454	0.463858
min	0.00000	0.000000	0.000000
25%	0.00000	0.000000	0.000000
50%	0.00000	0.000000	0.000000
75%	0.00000	2.000000	1.000000
max	6.00000	15.000000	1.000000

```
[6]: df.isnull().any().sum()
```

```
[6]: 2
```

```
[7]: df.isnull().sum()
```

```
[7]: ID                0
     AGE                0
     GENDER            0
     RACE               0
     DRIVING_EXPERIENCE 0
     EDUCATION          0
     INCOME             0
     CREDIT_SCORE       982
     VEHICLE_OWNERSHIP  0
     VEHICLE_YEAR       0
     MARRIED            0
     CHILDREN           0
     POSTAL_CODE        0
     ANNUAL_MILEAGE     957
     VEHICLE_TYPE       0
     SPEEDING_VIOLATIONS 0
     DUIS               0
     PAST_ACCIDENTS     0
     OUTCOME            0
     dtype: int64
```

```
[8]: df.shape
```

[8]: (10000, 19)

```
[9]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 19 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   ID                                    10000 non-null  int64
 1   AGE                                   10000 non-null  object
 2   GENDER                               10000 non-null  object
 3   RACE                                  10000 non-null  object
 4   DRIVING_EXPERIENCE                   10000 non-null  object
 5   EDUCATION                            10000 non-null  object
 6   INCOME                                10000 non-null  object
 7   CREDIT_SCORE                          9018 non-null   float64
 8   VEHICLE_OWNERSHIP                    10000 non-null  float64
 9   VEHICLE_YEAR                          10000 non-null  object
10   MARRIED                              10000 non-null  float64
11   CHILDREN                             10000 non-null  float64
12   POSTAL_CODE                          10000 non-null  int64
13   ANNUAL_MILEAGE                       9043 non-null   float64
14   VEHICLE_TYPE                          10000 non-null  object
15   SPEEDING_VIOLATIONS                  10000 non-null  int64
16   DUIS                                  10000 non-null  int64
17   PAST_ACCIDENTS                       10000 non-null  int64
18   OUTCOME                              10000 non-null  float64
dtypes: float64(6), int64(5), object(8)
memory usage: 1.4+ MB
```

```
[10]: df[df.isnull().any(axis=1)]
```

[illegible]

13	upper class	0.591260		1.0	before 2015	0.0
15	upper class	0.762798		0.0	after 2015	1.0
16	upper class	0.796175		1.0	before 2015	1.0
17	poverty	NaN		0.0	before 2015	1.0
18	upper class	0.680594		1.0	before 2015	0.0
...	...	...	...	...	...	...
9977	upper class	0.710640		1.0	after 2015	0.0
9981	working class	NaN		1.0	before 2015	0.0
9985	working class	NaN		1.0	before 2015	0.0
9988	poverty	NaN		0.0	before 2015	0.0
9996	middle class	0.522231		1.0	after 2015	0.0

	CHILDREN	POSTAL_CODE	ANNUAL_MILEAGE	VEHICLE_TYPE	SPEEDING_VIOLATIONS	\
13	1.0	10238	NaN	sedan		0
15	0.0	10238	NaN	sedan		0
16	1.0	32765	NaN	sedan		10
17	0.0	32765	12000.0	sedan		0
18	1.0	32765	NaN	sedan		0
...	...	...	...	...	...	...
9977	1.0	32765	NaN	sedan		0
9981	1.0	10238	11000.0	sedan		0
9985	1.0	10238	11000.0	sedan		0
9988	0.0	10238	NaN	sedan		1
9996	1.0	32765	NaN	sedan		1

	DUIS	PAST_ACCIDENTS	OUTCOME
13	0	0	0.0
15	0	0	0.0
16	2	1	0.0
17	0	0	1.0
18	0	0	1.0
...	...	...	...
9977	0	0	0.0
9981	0	0	0.0
9985	0	0	0.0
9988	0	2	1.0
9996	0	0	0.0

[1851 rows x 19 columns]

```
[11]: df[df["CREDIT_SCORE"].isnull()]
```

```
[11]:
```

	ID	AGE	GENDER	RACE	DRIVING_EXPERIENCE	EDUCATION	\
17	24851	16-25	male	majority	0-9y	none	
23	217	16-25	male	majority	0-9y	none	
37	511757	40-64	female	majority	10-19y	none	
38	429947	65+	male	majority	30y+	university	

47	921097	40-64	female	majority	20-29y	university
...	...	...	...	...	...	...
9952	870405	40-64	female	majority	10-19y	university
9967	27406	26-39	female	majority	10-19y	high school
9981	366048	26-39	male	majority	0-9y	high school
9985	595418	16-25	male	minority	0-9y	high school
9988	479789	26-39	male	majority	10-19y	high school


		INCOME	CREDIT_SCORE	VEHICLE_OWNERSHIP	VEHICLE_YEAR	MARRIED	\
17		poverty	NaN	0.0	before 2015	1.0	
23		poverty	NaN	0.0	before 2015	0.0	
37		middle class	NaN	1.0	before 2015	1.0	
38		upper class	NaN	0.0	after 2015	0.0	
47		upper class	NaN	1.0	after 2015	1.0	
...	...	...	...	...	...	...	...
9952		upper class	NaN	1.0	after 2015	1.0	
9967		middle class	NaN	0.0	before 2015	0.0	
9981		working class	NaN	1.0	before 2015	0.0	
9985		working class	NaN	1.0	before 2015	0.0	
9988		poverty	NaN	0.0	before 2015	0.0	

	CHILDREN	POSTAL_CODE	ANNUAL_MILEAGE	VEHICLE_TYPE	SPEEDING_VIOLATIONS	\
17	0.0	32765	12000.0	sedan		0
23	0.0	10238	17000.0	sedan		0
37	1.0	10238	11000.0	sedan		2
38	1.0	10238	12000.0	sports car		6
47	1.0	92101	11000.0	sedan		3
...	...	...	...	...	...	...
9952	1.0	32765	5000.0	sedan		1
9967	0.0	92101	13000.0	sedan		1
9981	1.0	10238	11000.0	sedan		0
9985	1.0	10238	11000.0	sedan		0
9988	0.0	10238	NaN	sedan		1

	DUIS	PAST_ACCIDENTS	OUTCOME
17	0	0	1.0
23	0	0	0.0
37	0	1	0.0
38	0	5	0.0
47	0	2	0.0
...	...	...	...
9952	0	0	0.0
9967	0	0	0.0
9981	0	0	0.0
9985	0	0	0.0
9988	0	2	1.0

[982 rows x 19 columns]

```
[12]: plt.figure(figsize=(8,5))
      df['CREDIT_SCORE'].plot(kind='kde')
      plt.show()
```



```
[13]: df["CREDIT_SCORE"].fillna(df["CREDIT_SCORE"].median(),inplace=True)
```

```
[14]: df["ANNUAL_MILEAGE"].fillna(df["ANNUAL_MILEAGE"].median(),inplace=True)
```

```
[15]: df.isna().sum()
```

```
[15]: ID          0
      AGE         0
      GENDER      0
      RACE        0
      DRIVING_EXPERIENCE  0
      EDUCATION   0
      INCOME      0
      CREDIT_SCORE  0
      VEHICLE_OWNERSHIP  0
      VEHICLE_YEAR  0
      MARRIED     0
      CHILDREN    0
      POSTAL_CODE  0
```

```

ANNUAL_MILEAGE      0
VEHICLE_TYPE        0
SPEEDING_VIOLATIONS 0
DUI                 0
PAST_ACCIDENTS      0
OUTCOME             0
dtype: int64

```

```
[16]: df.drop(["ID"],axis=1,inplace=True)
```

```
[17]: df.head()
```

```

[17]:      AGE  GENDER      RACE DRIVING_EXPERIENCE  EDUCATION      INCOME  \
0    65+  female  majority           0-9y  high school  upper class
1   16-25   male  majority           0-9y      none  poverty
2   16-25  female  majority           0-9y  high school  working class
3   16-25   male  majority           0-9y  university  working class
4   26-39   male  majority          10-19y      none  working class

      CREDIT_SCORE  VEHICLE_OWNERSHIP  VEHICLE_YEAR  MARRIED  CHILDREN  \
0      0.629027           1.0  after 2015      0.0      1.0
1      0.357757           0.0  before 2015      0.0      0.0
2      0.493146           1.0  before 2015      0.0      0.0
3      0.206013           1.0  before 2015      0.0      1.0
4      0.388366           1.0  before 2015      0.0      0.0

      POSTAL_CODE  ANNUAL_MILEAGE  VEHICLE_TYPE  SPEEDING_VIOLATIONS  DUI  \
0         10238         12000.0      sedan           0      0
1         10238         16000.0      sedan           0      0
2         10238         11000.0      sedan           0      0
3         32765         11000.0      sedan           0      0
4         32765         12000.0      sedan           2      0

      PAST_ACCIDENTS  OUTCOME
0              0      0.0
1              0      1.0
2              0      0.0
3              0      0.0
4              1      1.0

```

```
[18]: df["AGE"].unique()#discrete
```

```
[18]: array(['65+', '16-25', '26-39', '40-64'], dtype=object)
```

```
[19]: df["GENDER"].unique()#nominal
```

```
[19]: array(['female', 'male'], dtype=object)
```



```

[20]: df["RACE"].unique()#nominal

[20]: array(['majority', 'minority'], dtype=object)

[21]: df["DRIVING_EXPERIENCE"].unique()#discrete

[21]: array(['0-9y', '10-19y', '20-29y', '30y+'], dtype=object)

[22]: df["EDUCATION"].unique()#ordinal

[22]: array(['high school', 'none', 'university'], dtype=object)

[23]: df["INCOME"].unique()#ordinal

[23]: array(['upper class', 'poverty', 'working class', 'middle class'],
          dtype=object)

[24]: df["CREDIT_SCORE"].unique()#continous

[24]: array([0.62902731, 0.35775712, 0.49314579, ..., 0.47094023, 0.36418478,
          0.43522478])

[25]: df["VEHICLE_OWNERSHIP"].unique()#discrete

[25]: array([1., 0.])

[26]: df["VEHICLE_YEAR"].unique()#nominal

[26]: array(['after 2015', 'before 2015'], dtype=object)

[27]: df["MARRIED"].unique()#discrete

[27]: array([0., 1.])

[28]: df["CHILDREN"].unique()#discrete

[28]: array([1., 0.])

[29]: df["POSTAL_CODE"].unique()#nominal

[29]: array([10238, 32765, 92101, 21217], dtype=int64)

[30]: df["ANNUAL_MILEAGE"].unique()#discrete

[30]: array([12000., 16000., 11000., 13000., 14000., 10000., 8000., 18000.,
          17000., 7000., 15000., 9000., 5000., 6000., 19000., 4000.,
          3000., 2000., 20000., 21000., 22000.])

```

```
[31]: df["VEHICLE_TYPE"].unique()#nominal
```

```
[31]: array(['sedan', 'sports car'], dtype=object)
```

```
[32]: df["SPEEDING_VIOLATIONS"].unique()#continuous
```

```
[32]: array([ 0,  2,  3,  7,  6,  4, 10, 13,  1,  5,  9,  8, 12, 11, 15, 17, 19,  
        18, 16, 14, 22], dtype=int64)
```

```
[33]: df["DUI"].unique()#continuous
```


```
[33]: array([0, 2, 1, 3, 4, 5, 6], dtype=int64)
```

```
[34]: df["PAST_ACCIDENTS"].unique()#continuous
```

```
[34]: array([ 0,  1,  3,  7,  2,  5,  4,  6,  8, 10, 11,  9, 12, 14, 15],  
        dtype=int64)
```

```
[35]: df["OUTCOME"].value_counts().plot(kind='bar')#discrete
```


```
[35]: <AxesSubplot:>
```



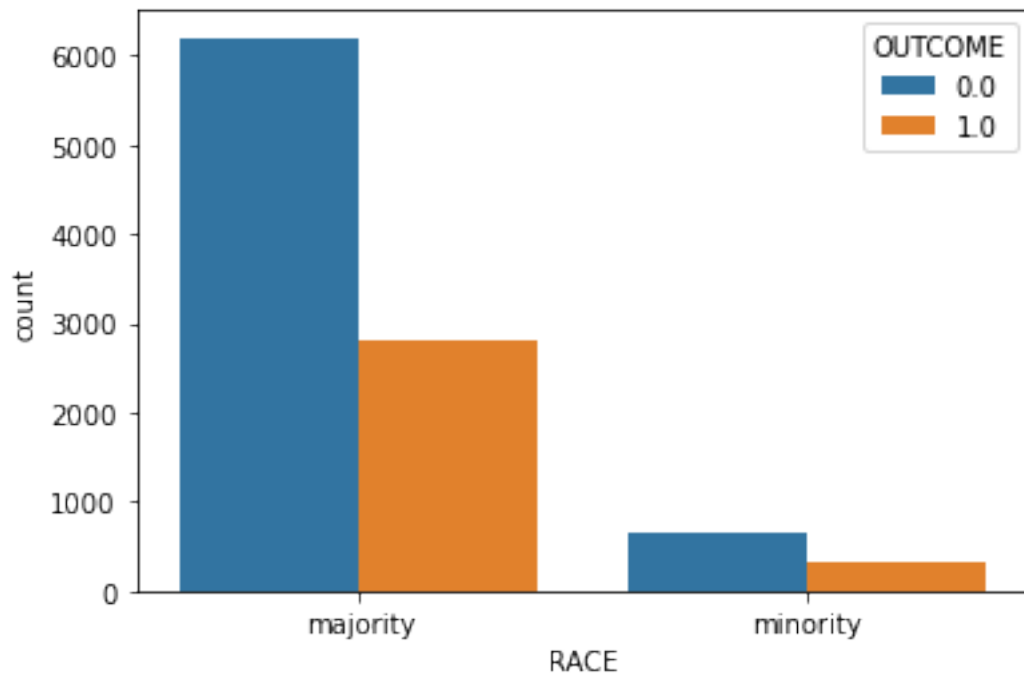
```
[36]: df["GENDER"].value_counts()
```

```
[36]: female    5010  
      male      4990  
      Name: GENDER, dtype: int64
```

```
[37]: sns.countplot(data=df,x="GENDER",hue="OUTCOME")  
      plt.show()
```




```
[38]: sns.countplot(data=df,x="RACE",hue="OUTCOME")  
      plt.show()
```



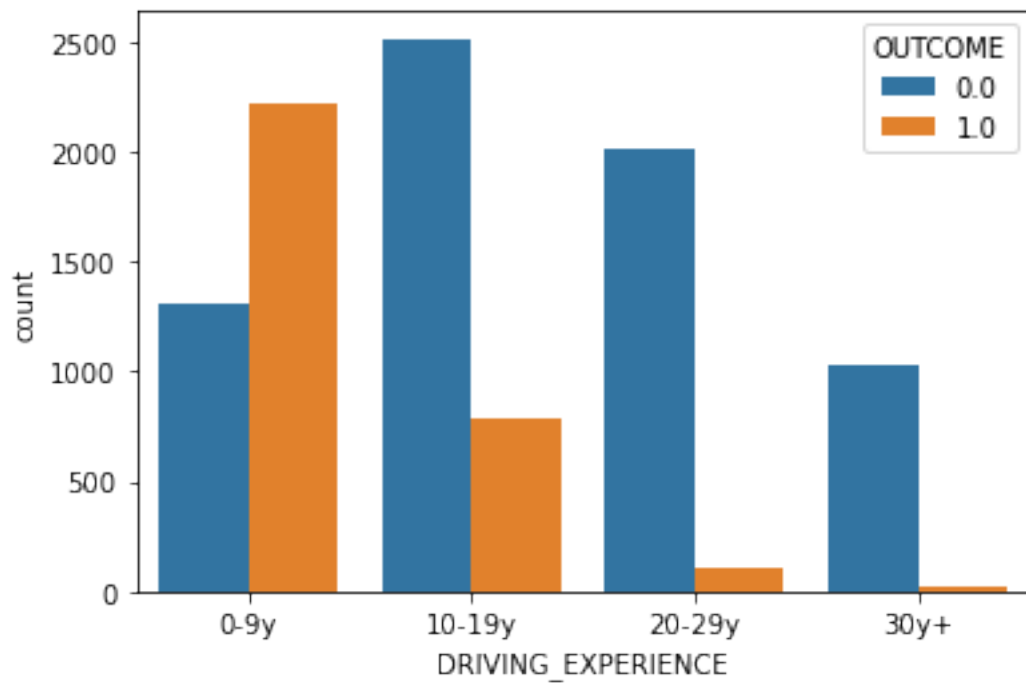
```
[39]: df["DRIVING_EXPERIENCE"].value_counts()
```

```
[39]: 0-9y      3530
      10-19y  3299
      20-29y  2119
      30y+   1052
      Name: DRIVING_EXPERIENCE, dtype: int64
```

```
[40]: sns.countplot(df["DRIVING_EXPERIENCE"])
      plt.show()
```




```
[41]: sns.countplot(data=df, x="DRIVING_EXPERIENCE", hue="OUTCOME")  
plt.show()
```



```
[42]: df["EDUCATION"].value_counts()
```

```
[42]: high school    4157  
      university    3928  
      none         1915  
      Name: EDUCATION, dtype: int64
```


```
[43]: sns.countplot(df["EDUCATION"])  
      plt.show()
```



```
[44]: df["INCOME"].value_counts()
```


```
[44]: upper class    4336  
      middle class  2138  
      poverty      1814  
      working class 1712  
      Name: INCOME, dtype: int64
```

```
[45]: sns.countplot(df["INCOME"])  
      plt.show()
```



```
[46]: df["OUTCOME"].value_counts().plot(kind='pie', autopct="%0.2f%%")
```


```
[46]: <AxesSubplot:ylabel='OUTCOME'>
```



```
[47]: df["VEHICLE_TYPE"].value_counts()
```

```
[47]: sedan          9523  
      sports car     477  
      Name: VEHICLE_TYPE, dtype: int64
```

```
[48]: sns.countplot(data=df,x='VEHICLE_TYPE',hue='OUTCOME')  
      plt.show()
```




```
[49]: df["DUI"].value_counts()
```

```
[49]: 0      8118  
      1      1470  
      2       331  
      3        68  
      4         10  
      5          2  
      6          1  
      Name: DUI, dtype: int64
```

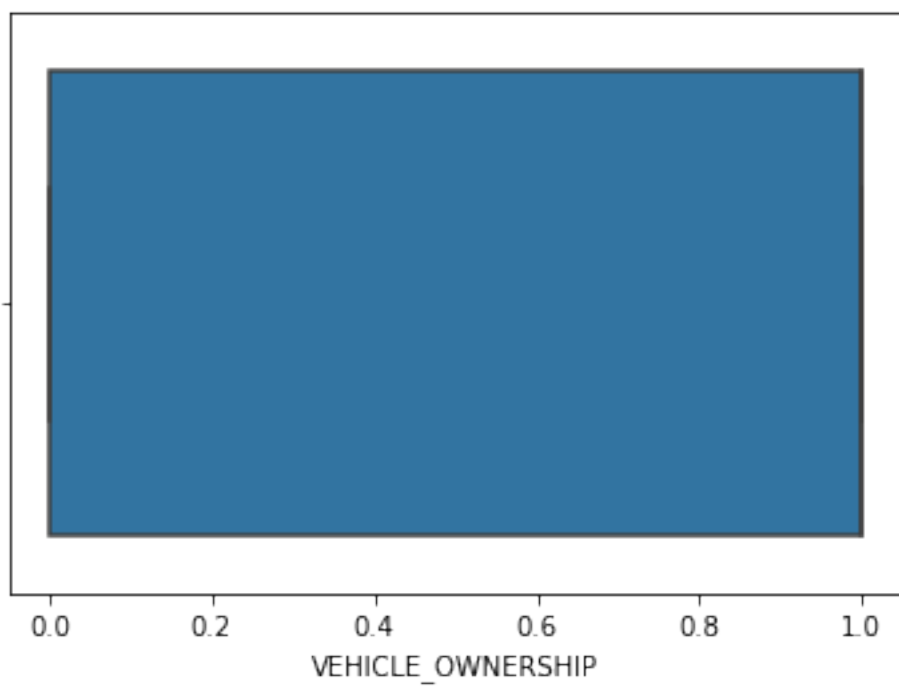

```
[50]: sns.countplot(df["DUI"])  
      plt.show()
```

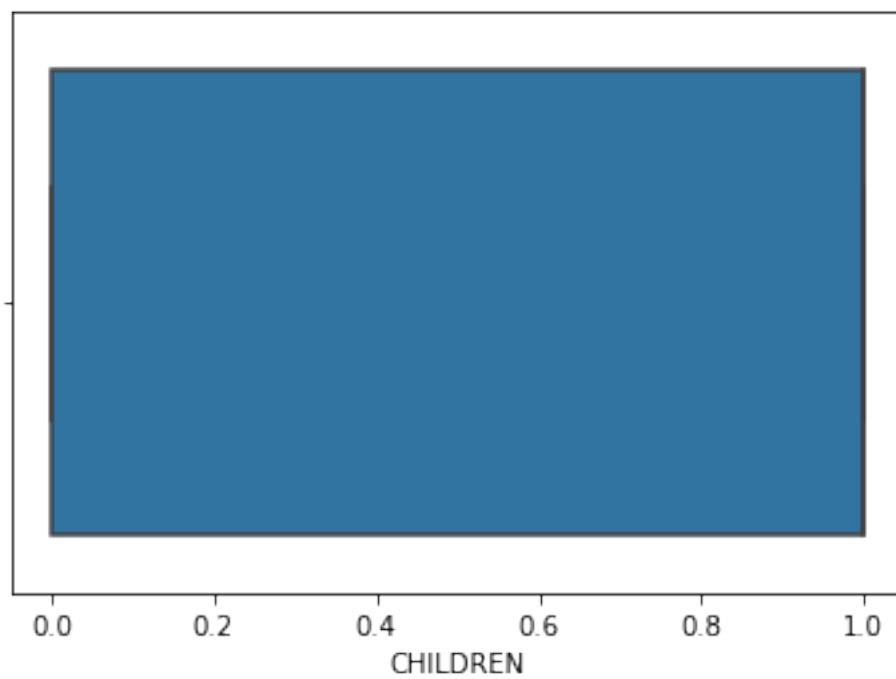
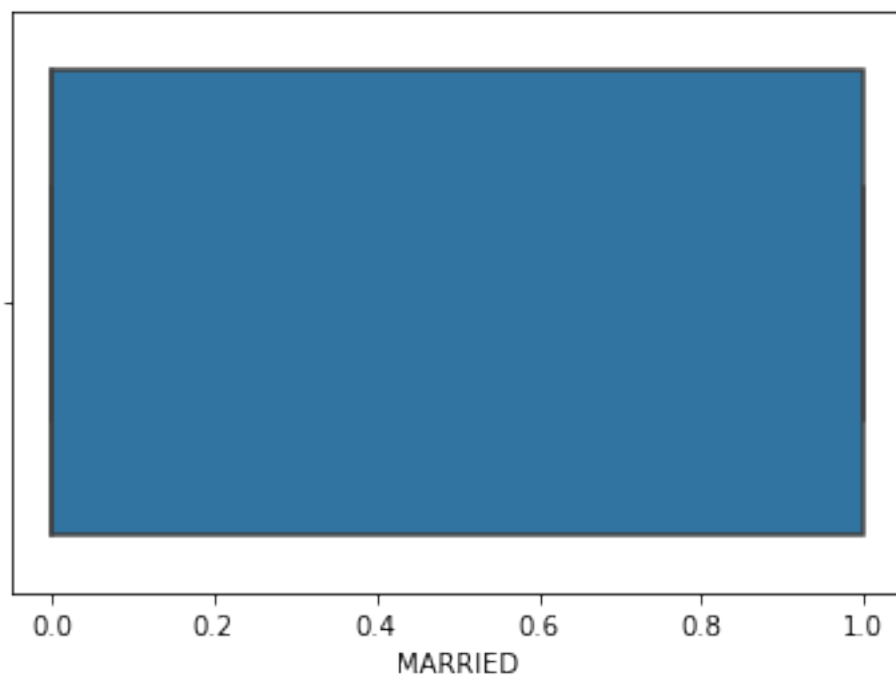







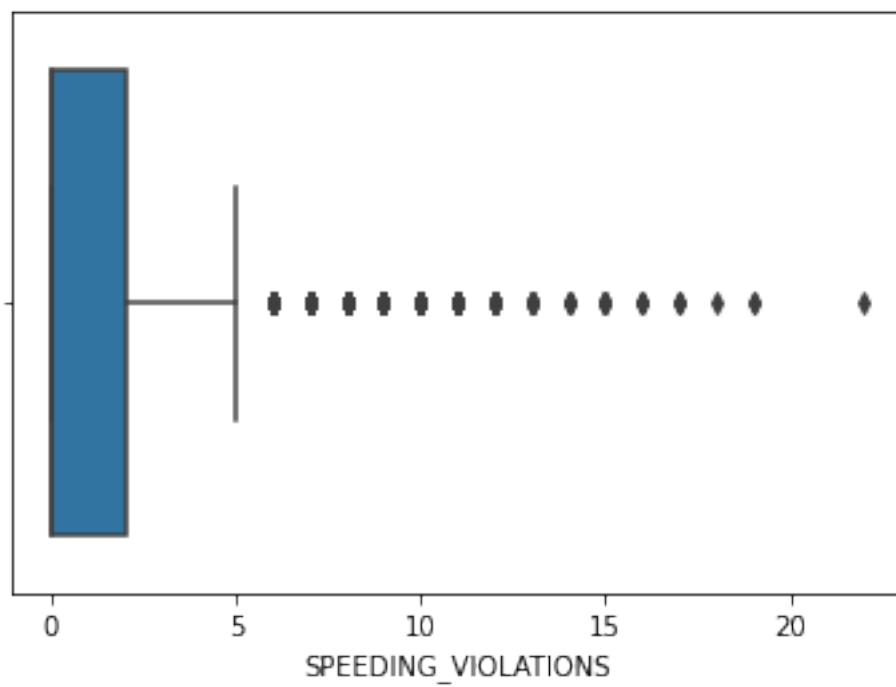
```
[51]: n=df.select_dtypes(exclude='object')
```

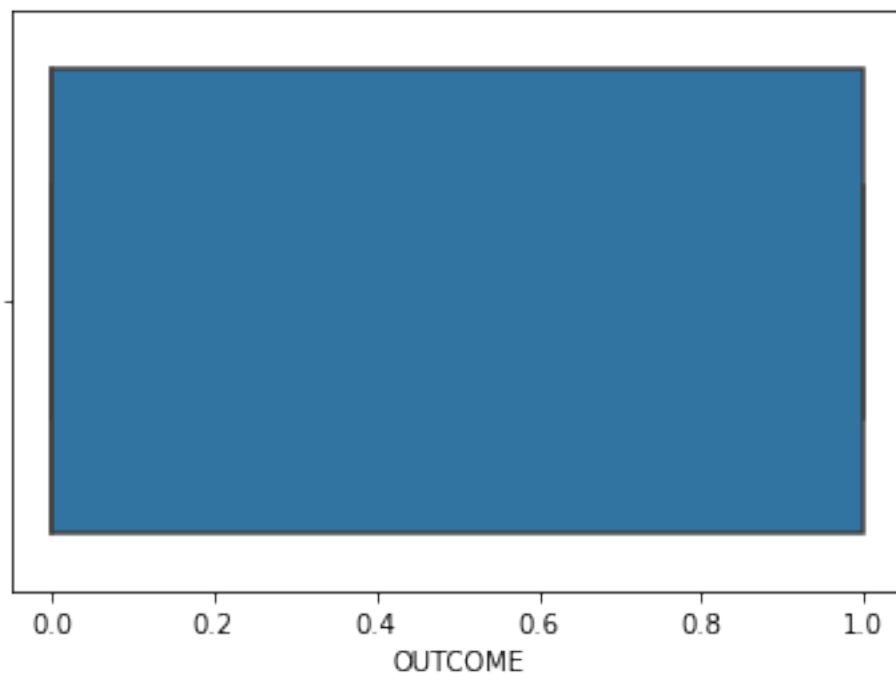
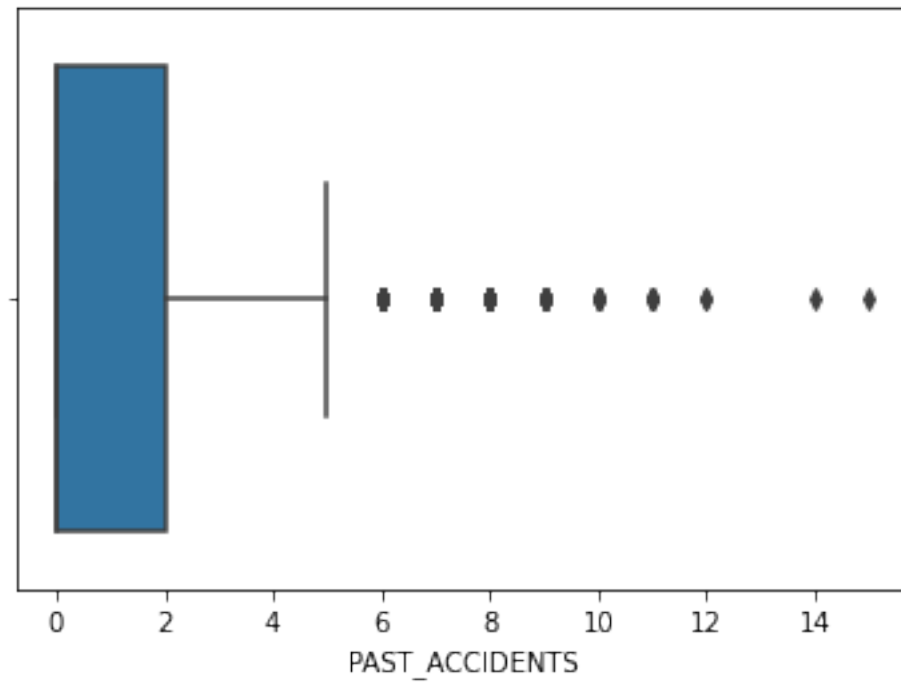
```
[52]: for i in n.columns:  
        sns.boxplot(data=n,x=i)  
        plt.show()
```











```
[53]: df.dtypes
```

```
[53]: AGE                object
      GENDER            object
      RACE              object
      DRIVING_EXPERIENCE object
      EDUCATION         object
      INCOME            object
      CREDIT_SCORE      float64
      VEHICLE_OWNERSHIP float64
      VEHICLE_YEAR      object
      MARRIED           float64
      CHILDREN          float64
      POSTAL_CODE       int64
      ANNUAL_MILEAGE    float64
      VEHICLE_TYPE      object
      SPEEDING_VIOLATIONS int64
      DUIS              int64
      PAST_ACCIDENTS    int64
      OUTCOME           float64
      dtype: object
```

```
[54]: from sklearn.preprocessing import LabelEncoder
      lr=LabelEncoder()
```

```
[55]: df['AGE']=lr.fit_transform(df['AGE'])
      df['GENDER']=lr.fit_transform(df['GENDER'])
      df['RACE']=lr.fit_transform(df['RACE'])
      df['DRIVING_EXPERIENCE']=lr.fit_transform(df['DRIVING_EXPERIENCE'])
      df['EDUCATION']=lr.fit_transform(df['EDUCATION'])
      df['INCOME']=lr.fit_transform(df['INCOME'])
      df['VEHICLE_YEAR']=lr.fit_transform(df['VEHICLE_YEAR'])
      df['VEHICLE_TYPE']=lr.fit_transform(df['VEHICLE_TYPE'])
```

```
[56]: df.dtypes
```

```
[56]: AGE                int32
      GENDER            int32
      RACE              int32
      DRIVING_EXPERIENCE int32
      EDUCATION         int32
      INCOME            int32
      CREDIT_SCORE      float64
      VEHICLE_OWNERSHIP float64
      VEHICLE_YEAR      int32
      MARRIED           float64
      CHILDREN          float64
      POSTAL_CODE       int64
      ANNUAL_MILEAGE    float64
```

```

VEHICLE_TYPE          int32
SPEEDING_VIOLATIONS   int64
DUI                   int64
PAST_ACCIDENTS        int64
OUTCOME               float64
dtype: object

```

```
[57]: df.head()
```

```

[57]:   AGE  GENDER  RACE  DRIVING_EXPERIENCE  EDUCATION  INCOME  CREDIT_SCORE  \
0     3      0     0          0          0      2      0.629027
1     0      1     0          0          1      1      0.357757
2     0      0     0          0          0      3      0.493146
3     0      1     0          0          2      3      0.206013
4     1      1     0          1          1      3      0.388366

      VEHICLE_OWNERSHIP  VEHICLE_YEAR  MARRIED  CHILDREN  POSTAL_CODE  \
0                   1.0            0      0.0      1.0      10238
1                   0.0            1      0.0      0.0      10238
2                   1.0            1      0.0      0.0      10238
3                   1.0            1      0.0      1.0      32765
4                   1.0            1      0.0      0.0      32765

      ANNUAL_MILEAGE  VEHICLE_TYPE  SPEEDING_VIOLATIONS  DUI  PAST_ACCIDENTS  \
0          12000.0          0          0          0          0
1          16000.0          0          0          0          0
2          11000.0          0          0          0          0
3          11000.0          0          0          0          0
4          12000.0          0          2          0          1

      OUTCOME
0          0.0
1          1.0
2          0.0
3          0.0
4          1.0

```

```
[58]: car_insurance=df.values
```

```
[59]: df.shape
```

```
[59]: (10000, 18)
```

```

[60]: x=car_insurance[:,0:17]
      y=car_insurance[:,17]

```



```
[61]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=0)
```

## 1 Logistic Regression

```
[62]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import
↳classification_report,confusion_matrix,accuracy_score
```

```
[63]: logistic_model = LogisticRegression().fit(x_train,y_train)
ypredicted = logistic_model.predict(x_test)
ypredicted
```

```
[63]: array([1., 0., 0., ..., 0., 0., 0.] )
```

## 2 Evaluation for Logistic Regression

```
[64]: Variance=np.var(ypredicted)
Variance
```

```
[64]: 0.20391822222222222
```

```
[65]: SE=np.mean((np.mean(ypredicted)-y)**2)
bias=SE-Variance
bias
```

```
[65]: 0.012007022222222219
```

```
[66]: print("Confusion Matrix")
matrix = confusion_matrix(y_test,ypredicted)
print(matrix)
```

```
Confusion Matrix
[[1842  203]
 [ 302  653]]
```

```
[67]: print("\nClassification Report")
report = classification_report(y_test,ypredicted)
print(report)
```

```
Classification Report
              precision    recall  f1-score   support
```

0.0	0.86	0.90	0.88	2045
1.0	0.76	0.68	0.72	955
accuracy			0.83	3000
macro avg	0.81	0.79	0.80	3000
weighted avg	0.83	0.83	0.83	3000

```
[68]: lr_accuracy = accuracy_score(y_test, ypredicted)
lr_accuracy
print('Logistic Regression Accuracy of Scikit Model: {:.2f}%'.
      ↪format(lr_accuracy*100))
```

Logistic Regression Accuracy of Scikit Model: 83.17%

```
[69]: from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import RepeatedStratifiedKFold
```

```
[70]: solvers = ['newton-cg', 'lbfgs', 'liblinear']
penalty = ['l2']
# define models and parameters
model = LogisticRegression()
c_values = [100, 1000, 1.0, 0.1, 0.01]
# define grid search
grid = dict(solver=solvers,penalty=penalty,C=c_values)
cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
grid_search = GridSearchCV(estimator=model, param_grid=grid, n_jobs=-1, cv=cv,
↪scoring='accuracy',error_score=0)
grid_result = grid_search.fit(x_train,y_train)
# summarize results
print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
means = grid_result.cv_results_['mean_test_score']
stds = grid_result.cv_results_['std_test_score']
params = grid_result.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean*100, stdev*100, param))
```

```
Best: 0.845333 using {'C': 100, 'penalty': 'l2', 'solver': 'newton-cg'}
84.533333 (1.070275) with: {'C': 100, 'penalty': 'l2', 'solver': 'newton-cg'}
80.647619 (2.359734) with: {'C': 100, 'penalty': 'l2', 'solver': 'lbfgs'}
80.528571 (1.225661) with: {'C': 100, 'penalty': 'l2', 'solver': 'liblinear'}
84.519048 (1.053974) with: {'C': 1000, 'penalty': 'l2', 'solver': 'newton-cg'}
80.819048 (2.513618) with: {'C': 1000, 'penalty': 'l2', 'solver': 'lbfgs'}
80.528571 (1.225661) with: {'C': 1000, 'penalty': 'l2', 'solver': 'liblinear'}
84.509524 (0.994326) with: {'C': 1.0, 'penalty': 'l2', 'solver': 'newton-cg'}
80.833333 (2.426914) with: {'C': 1.0, 'penalty': 'l2', 'solver': 'lbfgs'}
80.528571 (1.225661) with: {'C': 1.0, 'penalty': 'l2', 'solver': 'liblinear'}
84.390476 (0.895137) with: {'C': 0.1, 'penalty': 'l2', 'solver': 'newton-cg'}
```

```
81.076190 (2.496728) with: {'C': 0.1, 'penalty': 'l2', 'solver': 'lbfgs'}
80.528571 (1.233407) with: {'C': 0.1, 'penalty': 'l2', 'solver': 'liblinear'}
83.533333 (0.972362) with: {'C': 0.01, 'penalty': 'l2', 'solver': 'newton-cg'}
80.600000 (1.885642) with: {'C': 0.01, 'penalty': 'l2', 'solver': 'lbfgs'}
80.509524 (1.139688) with: {'C': 0.01, 'penalty': 'l2', 'solver': 'liblinear'}
```

```
[71]: grid_result.score(x_train,y_train)*100
```

```
[71]: 84.81428571428572
```

```
[73]: from sklearn.ensemble import RandomForestClassifier
      from sklearn.neighbors import KNeighborsClassifier
      from sklearn.tree import DecisionTreeClassifier
      from sklearn.svm import SVC
      from sklearn.naive_bayes import GaussianNB
      from sklearn.model_selection import cross_val_score
```

```
[75]: # Spot-Check Algorithms
      models = []
      models.append(('RF', RandomForestClassifier()))
      models.append(('GNB', GaussianNB()))
      models.append(('KNN', KNeighborsClassifier()))
      models.append(('DSC', DecisionTreeClassifier(random_state = 1, max_depth=2)))
      models.append(('SVM', SVC()))
      # evaluate each model in turn
      results = []
      names = []
      for name, model in models:
          kfold = RepeatedStratifiedKFold(n_splits=10, n_repeats = 3, random_state=1)
          cv_results = cross_val_score(model,x_train,y_train, cv=kfold,
          ↪scoring='accuracy')
          results.append(cv_results)
          names.append(name)
          msg = "%s: %.2f (%.3f)" % (name, cv_results.mean(), cv_results.std())
          print(msg)
```

```
RF: 0.84 (0.011)
GNB: 0.73 (0.017)
KNN: 0.79 (0.014)
DSC: 0.81 (0.015)
SVM: 0.70 (0.004)
```

```
[76]: # evaluate model 1
      model1 = RandomForestClassifier()
      cv1 = RepeatedStratifiedKFold(n_splits = 10, n_repeats = 3, random_state = 1)
      scores1 = cross_val_score(model1,x_train,y_train, scoring = 'accuracy', cv =
      ↪cv1, n_jobs = -1)
```


```

print('RFC Mean Accuracy: %.1f%% +/- (0.3f)' % (np.mean(scores1*100), np.
    ↳std(scores1)))
# evaluate model 2
model2 = DecisionTreeClassifier(random_state = 1, max_depth=2)
cv2 = RepeatedStratifiedKFold(n_splits = 10, n_repeats = 3, random_state = 1)
scores3 = cross_val_score(model2,x_train,y_train, scoring = 'accuracy', cv =
    ↳cv2, n_jobs = -1)
print('DecisionTreeClassifier Mean Accuracy: %.1f%% +/- (0.3f)' % (np.
    ↳mean(scores3*100), np.std(scores3)))
# plot the results
plt.boxplot([scores1, scores3], labels=['RFC', 'DTC'], showmeans=True)
plt.show()

```

RFC Mean Accuracy: 83.7% +/- (0.012)

DecisionTreeClassifier Mean Accuracy: 80.9% +/- (0.015)



```

[77]: from mlxtend.evaluate import paired_ttest_5x2cv
# check if difference between algorithms is real
t, p = paired_ttest_5x2cv(estimator1=model1,
    estimator2=model2,
    X=x,
    y=y,
    scoring='accuracy',
    random_seed=1)
# summarize

```

```

print(f'The P-value is = {p:.3f}')
print(f'The t-statistics is = {t:.3f}')
# interpret the result
if p <= 0.05:
    print('Since  $p < 0.05$ , We can reject the null-hypothesis that both models
    ↪ perform equally well on this dataset. We may conclude that the two
    ↪ algorithms are significantly different.')
else:
    print('Since  $p > 0.05$ , we cannot reject the null hypothesis and may conclude
    ↪ that the performance of the two algorithms is not significantly different.')

```

The P-value is = 0.021

The t-statistics is = 3.339

Since  $p < 0.05$ , We can reject the null-hypothesis that both models perform equally well on this dataset. We may conclude that the two algorithms are significantly different.

[ ]: